Logical database design

- The learning objectives for this week are:
    - Knowing what is the objective **logical database design**
    - Knowing how to **derive relations** from entity-relationship diagrams
   - Knowing how to determine **natural primary keys** and **foreign keys**

Logical database design

- The typical main phases in a database design process are:
  1. Conceptual database design
  2. Logical database design
  3. Physical database design
- Each phase from top to bottom adds more detail to the design
- We have familiarized ourselves with the **conceptual database design** by defining entities and their attributes and relationships based on the requirements
- The **logical database design** is the process of refining and translating the conceptual schema into a **logical database schema** based on a specific data model, for example the relational model

Conceptual vs. logical database schema

- Here's a comparison of conceptual (left) and logical (right) database schema of a course enrollment database:
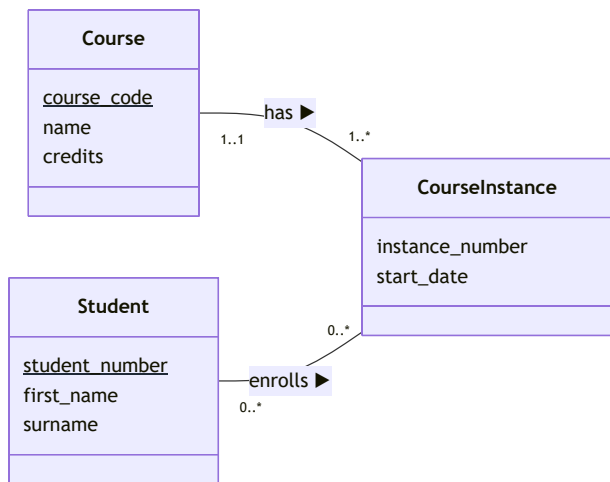

```
Course(course_code, name, credits)
Student(student_number, first_name, surname)

CourseInstance(course_code, instance_number, start_date)
    FOREIGN KEY (course_code) REFERENCES Course(course_code)

Enrollment(course_code, instance_number, student_number)
    FOREIGN KEY (course) REFERENCES Course(course_code)
    FOREIGN KEY (course_code, instance_number)
    REFERENCES CourseInstance(course_code, instance_number)
    FOREIGN KEY (student_number)
    REFERENCES Student(student_number)
```

Logical database design

- Entity types, attributes, and relationship types can be directly transformed into relations with some simple rules
- Typically, the logical database design process includes the following types of activities:
  1. Deriving relations for the logical data model
  2. Validating relations using normalisation
  3. Validating relations against user transaction
  4. Double-checking integrity constraints
  5. Reviewing logical data model with user

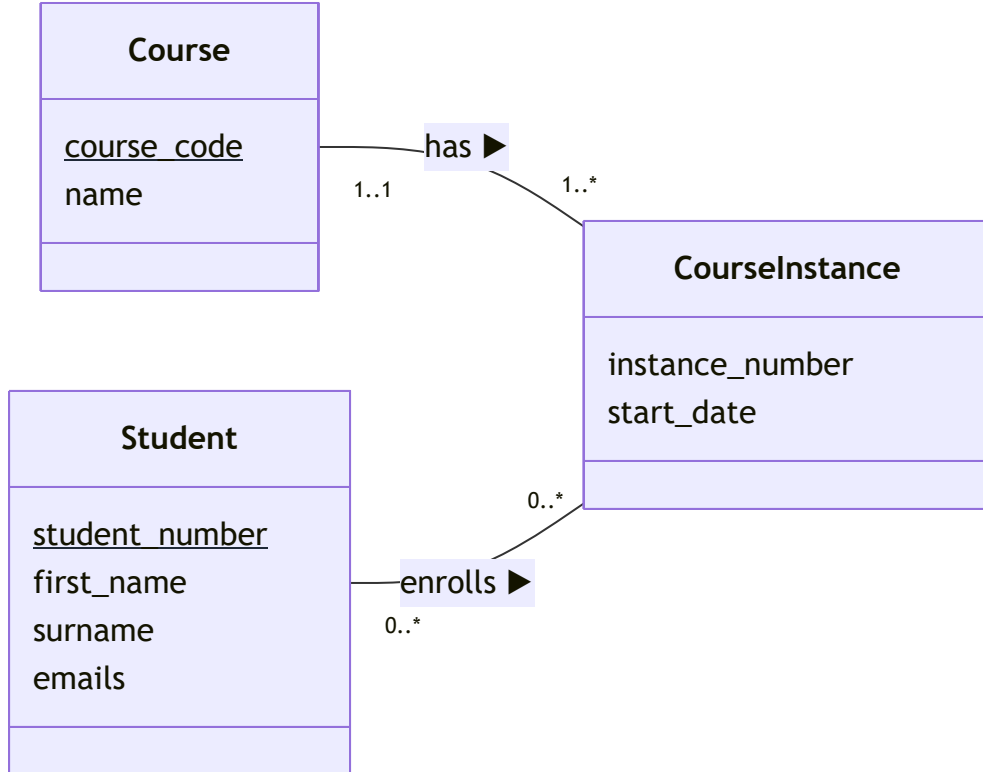# Deriving relations for the logical data model

- The process starts by **deriving relations for the logical data model**, which includes:

  1. Creating the relations

  2. Refining the attributes

  3. Determining primary and foreign keys

  4. Determining other types of integrity constraints

Creating relations

- To transform a conceptual schema into a relational database schema, we might need **additional relations** to establish specific relationships
- We create the relations in the following manner:
  - For **each entity type**, we create a relation that includes all simple (single-value) attributes of the entity
  - For **many-to-many relationship types** ("...*" on both sides of the relationship), we create a **"bridge relation"** to represent the relationship
  - For **multi-valued attributes**, we create a new relation to represent the multi-valued attribute. For example, a person may have several phone numbers, but multi-valued attributes are not allow in relations

Example of creating relations

- Let's consider creating relations for the following conceptual model:

Course

course_code
name

has ▶

1..1   1..*

CourseInstance

instance_number
start_date

Student

student_number
first_name
surname
emails

enrolls ▶

0..*

0..*

Example of creating relations

- In the example we have one **many-to-many relationship** between the *Student* and *CourseInstance* entities ("…*" on both sides of the "enrolls" relationship):

  - *"Student enrolls to zero or many course instances and course instance has zero or many students"*

  - In this case we create an additional **bridge relation**, for example *Enrollment*

- There is a **multi-valued attribute** *emails* in the *Student* entity

  - In this case we create an additional entity, for example *StudentEmail*

- This leaves with the following relations: *Student*, *CourseInstance*, *Enrollment* and *StudentEmail*

Refining the attributes

- Once we have created the relations, we need to refine the attributes in the following manner:
    - We divide a non-atomic attribute into smaller (atomic) attributes. For example student's home address can be divided into, city, postal code and street_address attributes
    - We define general (not DBMS-specific) attribute data-types, for example "string" or "integer"
    - We define which attributes can have `NULL` values. We should allow `NULL` in an attribute only based on **strong arguments**

Determining primary keys

```
Student(student_number, first_name, surname)
```

- There should be **exactly one primary key in each relation**
- The primary key can be either a **simple key** (single column, like `course_code` in the `Course` relation) or a **composite key** (several columns, like `course_code` and `instance_number` in the `CourseInstance` relation)
- By definition, the primary key should always satisfy the properties of **requiredness** (not `NULL`), **uniqueness** and **minimality**
- The primary key **should remain stable**. That is, primary key values should not be updated later
- The primary key should be **reasonably short**
- The primary key should have **no privacy issues**. For example social security number has privacy issues

Determining primary key for a weak entity type

- A **weak entity type** is an entity type that is dependent on the existence of another entity type

- For example *CourseGrade* is existence-dependent on *Student* and *CourseInstance*

- When a relation derived from a weak entity type, the natural primary key is partially or fully derived from the weak entity type's owner entity type

- For example, the natural primary key of the *CourseGrade* relation is a composite key that **includes columns from two foreign keys**, one referencing *Student* and other referencing *CourseInstance*

- The primary key of the weak entity's relation cannot be made until after the foreign keys have been determined for the relation

Surrogate keys

- If there is initially no suitable candidate key for a relation, then we cannot determine a natural primary key
- We have to take care of the situation by including an extra attribute in the relation to act as the primary key
- This kind of primary key is a **surrogate key**
- Surrogate keys are commonly generated values, such as incrementing or random numbers, like the `messageid` primary key in the data below

messageid	from	to	title	body
1	kalle.ilves@haaga-helia.fi	john.doe@gmail.com	Greeting	Hello John!
2	john.doe@gmail.com	kalle.ilves@haaga-helia.fi	Response	Hello Kalle!
…	…	…	…	…

Alternate keys

- Candidate keys that are not selected to be primary the key are called **alternate keys**

- We should consider the use of the **unique constraint** on alternate keys to make sure that their values remain unique:

```
Student (studentnumber, ssn, familyname, givenname)
    UNIQUE (ssn)
```

- Especially, when we are using a surrogate primary key, a unique constraint on at least one natural alternate key improves data quality
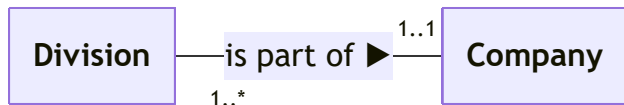
Determining foreign keys

- In a relational database, **relationships** are represented by the **primary key/foreign key mechanism**

- To know in which relation be need to place the foreign key, we need to identify the **relationship type** between the two entities

- During the course, we will use the following syntax to define the relation's foreign keys in the **relation schema**:

```
Course(course_code, name, credits)

CourseInstance(course_code, instance_number, start_date)
  FOREIGN KEY (course_code) REFERENCES Course(course_code)
```
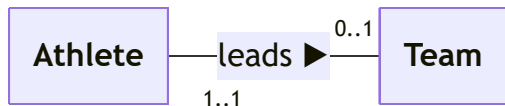
Determing foreign keys

- Most often the relationship type between two entities falls in **many-to-one**, **many-to-many** and **one-to-one** categories

- Example of **many-to-one relationship** with "…*" on one side and "…1" on another:

| Division | —is part of ▶— | Company |

Division —is part of ▶— Company (1..1 on Company side, 1..* on Division side)

- Example of **many-to-many relationship** with "…*" on both sides:

Athlete —participates ▶— Race (0..* on Race side, 0..* on Athlete side)

- Example of **one-to-one relationship** with "…1" on both sides:

Athlete —leads ▶— Team (0..1 on Team side, 1..1 on Athlete side)

Many-to-one relationship (1:N)

> *"Division has one or many companies"*
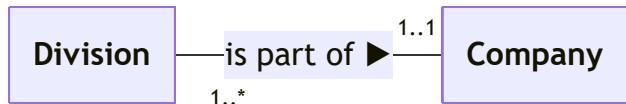
- Place a copy of the **parent relation's** ("…1" side) primary key into the **child relation** ("…*" side), to act as a foreign key

- If the child relation is derived from a weak entity type, then the primary key of the child relation is typically a composite key

```
Company (company_id, name)

Division (division_id, company_id, name)
  FOREIGN KEY (company_id) REFERENCES Company(company_id)
```

Many-to-many relationship (M:N)

"Athlete participates in zero or many races and race zero or many athlete participants"


- Create a **bridge relation** to represent the relationship and place a copy of the primary key from each of the parent relations into the bridge relation to act as foreign keys

- Typically, the bridge relation's primary key is a composite key that includes the both foreign keys

```
Athlete (athlete_id, first_name, family_name)
Race (race_id, name, date)
RaceParticipation(athlete_id, race_id)
  FOREIGN KEY (athlete_id) REFERENCES Athlete(athlete_id),
  FOREIGN KEY (race_id) REFERENCES Race(race_id)
```

One-to-one relationship (1:1)

> "Team has exactly one athlete leader"
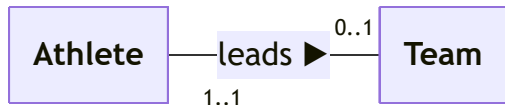
- In case of **mandatory participation** ("1...1" on one side only), place a copy of the primary key from the relation on the "1...1" side into the relation on the "0...1" side to act as the foreign key

- In case of **mandatory participation on both sides** we can usually **combine the two** into one relation

- In case **optional participation** ("0...1" on both sides) the foreign key can be placed in either relation

```
Athlete (athlete_id, first_name, family_name)
Team (team_id, athlete_id, name)
  FOREIGN KEY (athlete_id) REFERENCES Athlete(athlete_id)
```

Multi-value attributes

Employee
empno
family_name
given_name
phone_numbers

- A relation can't have attributes with **multiple values**, such as the *phone_numbers* attribute of the *Employee* entity type in this example (employee has many phone numbers)
- In such case, we must create a **new relation** to represent the multi-valued attribute, for example *EmployeePhone*
- We move the attribute from the original relation and place it to the new relation and place a copy of the parent relation's primary key into the child relation, to act as the foreign key

Multi-value attributes

Employee
empno
family_name
given_name
phone_numbers

- In this example, we would get the following relation schema:

```
Employee (empno, first_name, family_name)
EmployeePhone (phone_number, empno)
  FOREIGN KEY (empno) REFERENCES Employee(empno)
```

Summary

- The objective of logical database design is to translate the conceptual schema into a **logical database schema** based on a specific data model

- During the course, we represent logical database schema as **relation schema**, containing relations their attributes, primary key and foreign keys

- When we derive relations from entity types, we create a relation for each entity type

- Many-to-many relationship requires an additional **bridge relation**

- A **multi-valued attributes** requires an additional relation

- There should be exactly one primary key in each relation

- The foreign key placement depends on the **relationship type** (many-to-one, many-to-many or one-to-one)