Database normalisation

- The learning objectives for this week are:
    - Knowing the purpose of *database normalisation*
    - Knowing what is a *functional dependency*, a *partial dependency* and a *transitive dependency*
    - Knowing how to identify functional dependencies in a relation or table
    - Knowing the different *normal form* rules
    - Knowing how to formally check if a relation is in the *Boyce-Codd normal form* (BCNF)
    - Knowing how to *decompose a relation* into smaller relations if it is not in BCNF

Database normalisation

- *Database normalisation* is a formal technique of organizing data in a database in a way that *redundancy* and *incosistency* within the data is eliminated
- The objective of database normalisation is to ensure that:
  - Attributes with a *close logical relationship* (functional dependency) are found in the *same relation*
  - The relations do not display *hidden data redundancy*, which can cause update anomalies that violate database integrity
- The technique involves a set of *normalisation rules* that are defined as *normal forms*

Redundancy example

Let's consider *redundancy problems* with the following `Course_Enrolment` relation rows:

courseno	studentno	phone	enrolment_date
C001	10	1234	2025-04-01
C001	20	5555	2025-04-02
C002	30	8765	2025-04-01
C002	40	1414	2025-04-03
C002	10	1234	2025-04-07

Redundancy example

The student `10` *phone number is duplicated* causing redundancy in the data. While updating a phone number or inserting a new row, there's a risk of having multiple different phone numbers for the same student (inconsistency):

courseno	studentno	phone	enrolment_date
C001	10	1234 ⚠️	2025-04-01
...	...	...	...
C002	10	3338 ⚠️	2025-04-07

Database normalisation

- In a case of fixing an identified structural problem, normalisation involves *decomposing a relation into less redundant (and smaller) relations* without losing information

- When an *ER model is well designed*, the resulting correctly derived relations won't normally have such structural problems and they will meet the criteria of database normalisation

- Normalisation of candidate relations derived from ER diagrams is accomplished by analysing the *functional dependencies* (FDs) associated with those relations

Functional dependency

- *Functional dependency* (FD) describes the *relationship between attributes* in a relation

- With functional dependencies, we are interested in properties of the data that are true for *all the time*

- For example, if the *student number is unique*, the following property is true all the time:

  > The surname for a student whose student number is "a12345" is "Smith"

- So, *all the time* it is true that there is only one surname for each student

- By contrast, the following property might to be true for a sample set of students, but it is not true for all the time:

  > There is exactly one student whose surname is "Smith"

Functional dependency

- A functional dependency occurs when attribute A in a relation *uniquely determines* attribute B

- In other words: for each value of A there is *exactly one value* of B and that *holds all the time*. This can be written as `A → B`

- The *determinant* of a functional dependency refers to the attribute, or group of attributes, on the *left-hand side* of the arrow. In `A → B`, A is the determinant of B.

- On the *right-hand side*, there's the *dependent*. In `A → B`, B is the dependent of A.

Example of functional dependency

- Let's suppose that each student has a unique student number. In the relation below, *studentnumber* uniquely determines *surname* and *firstname*. That is, *studentnumber is the determinant of surname and firstname*:

```
Student (studentnumber, surname, firstname)
```

- In this example, there are the following two functional dependencies:

    - `studentnumber → surname`

    - `studentnumber → firstname`

Example of functional dependency

- Let's suppose the following table occurrence:

studentnumber	surname	firstname
a12345	Smith	John
a14444	Smith	Susan
a15555	Jones	Susan

- The *functional dependency* `studentnumber → surname` guarantees that the query below (that uses an existing student number) returns exactly one surname and that holds all the time:

```
SELECT surname FROM Student WHERE studentnumber = 'a12345'
```

Example of functional dependency

- `{A, B} → C` means that *A and B together uniquely determine* C. For example, `{course_code, implementation_number} → start_date`

- `A → B, C, D` means that *A uniquely determines B, C, and D*, For example, `course_code → course_name, language, credits`

Identifying undesired data redundancy

- Relations that *do not have* undesired data redundancy, *each determinant is a candidate key* (an unique attribute that is suitable for being the primary key)

- In such case *all arrows are arrows out of whole candidate keys* (simple or composite key)

- Let's consider the following relation *without data redundancy*:

  ```
  CourseOffering (coursecode, offeringnumber, startdate, teachernumber)
  ```

- In this relations there's for example the following functional dependency:
  - ✅ `{coursecode, offeringnumber} → startdate, teachernumber`

Identifying undesired data redundancy

- Relations that *have* undesired data redundancy, *there is a determinant that is not a candidate key*

- In such case *there is on arrow that is not an arrow out of a whole candidate key*

- Let's consider the following relation *with data redundancy*:

```
CourseOffering (coursecode, offeringnumber, coursename,
                startdate, teachernumber, surname)
```

Identifying undesired data redundancy

- In this relations there's for example the following functional depedencies:
    - ✅ `{coursecode, offeringnumber} → coursename, startdate, teachernumber, surname`
    - ❌ `coursecode → coursename`
    - ❌ `teachernumber → surname`
- In functional dependencies `coursecode → coursename` and `teachernumber → surname`, the determinants are not candidate keys
- With such functional dependencies, the relation has redundant data
- For example the teacher's surname is repeated unnecessarily, which can cause consistency issues for example when a teacher's surname is updated
- Instead, the teacher's information should be in a *separate relation*

Calculated attributes

- We *should not include* attributes in a relation that we can *derive* from other relations or *calculate*

- For example, let's suppose that the firm's total budget is the total of department budgets

- Therefore, *totalbudget* is a calculated attribute in the *Firm* relation

- The value of *totalbudget* should change whenever any department budget is changed in the firm

- From the data redundancy and integrity viewpoint, we have a problem here because total budget exists twice in the design:

```
Firm (firmno, firmname, totalbudget ❌)
Depertmant (deptno, deptname, deptbudget, firmno)
    FK (firmno) REFERENCES Firm (firmno)
```

Calculated attributes

- We shouldn't have the *totalbudget* attribute in the *Firm* relation, instead we can calculate it with the following query:

```
SELECT SUM(deptbudget) as totalbudget FROM Department
WHERE firmno = 'a1122'
```

Different kind of functional dependencies

- Functional dependencies can be categorized in the following categories:
    - *Non-trivial* and *trivial* functional dependencies
    - *Partial* and *full* functional dependencies
    - *Transitive* and *non-transitive* functional dependencies

Non-trivial and trivial functional dependencies

- `A → B` is *trivial functional dependency* if B is a subset of A

- `A → B` is *non-trivial functional dependency* if B is not a subset of A

- Let's consider the *CourseOffering* relations:

  ```
  CourseOffering (coursecode, offeringno, startdate)
  ```

- In the relation, `{coursecode, offeringno} → startdate` is a *non-trivial functional dependency*, because `startdate` is not a subset of `{coursecode, offeringno}`

- These, on the other are *trivial functional dependencies* of the relation:

  - `{coursecode, offeringno} → coursecode`

  - `{coursecode, offeringno} → {coursecode, offeringno}`

- In normalisation considerations we are only focusing on *non-trivial functional dependencies*

Normal forms

- *Normal form* refers to a set of normalisation rules that a database relation should follow inorder to be considered "normalized" and thus *well-organized*

- During the course we will cover the most common normal forms: *first normal form* (1NF), *second normal form* (2NF), *third normal form* (3NF) and *Boyce-Codd normal form* (BCNF)

- Each normal form from 1NF to BNCF adds more rules to the previous normal form

- For example, the 2NF includes all rules of the 1NF and additional rules

- The Boyce-Codd Normal Form (BCNF) is the strictest of these normal forms

First normal form (1NF)

- A relation is in the *first normal form* (1NF) if the following rules apply:
  - All attributes in a relation *must have atomic values*. No multi-valued attributes are allowed
  - A relation *must have a primary key* and all its *attributes must be dependent on the primary key*

Second normal form (2NF)

- A relation is in the *second normal form* (2NF) if the following rules apply:
  - Relation is in 1NF
  - Relation has no *partial functional dependencies*, meaning that there is no *part of a candidate key* that uniquely determines a *non-candidate-key* attribute
- Let's consider the following relation:

```
ClubMembership (empno, clubno, clubname, joindate)
```

- The relation has a *partial functional dependency* `{empno, clubno} → clubname`, because the functional dependency `clubno → clubname` exists in the relation
- That is, the relation *is not in 2NF*

Third normal form

- A relation is in the *third normal form* (3NF) if the following rules apply:
  - Relation is in 1NF

  - Relation has no functional dependency between two *non-candidate-key* attributes, meaning no *non-candidate-key* attribute is allowed to be *transitively* dependent on any *candidate key* within the relation

- Let's consider the following relation schema:

  ```
  Employee (empno, surname, firstname, deptno, deptname)
  ```

- The relation has a transitive functional dependency `deptno → deptname`, causing `deptname` to be *transitively dependent* on `empno` via `deptno`

- That is, the relation *is not in 3NF*

Boyce-Codd Normal Form (BCNF)

- We simplify the rules of BCNF we will have the following limitations during the course:
    - We only focusing on *non-trivial functional depdencies*
    - Instead of including any superkeys in our analysis, we narrow the analysis to *candidate keys*
    - We do not allow any attribute that *does not have a determinant* within the relation
- With these limitations the BCNF has the following rules for a relation:
    - Each determinant is a candidate key
    - All attribute values are atomic (single values)
    - There is a determinant for each attribute that is not contained in a candidate key

Boyce-Codd Normal Form (BCNF)

- Let's consider the following relation:

  ```
  Teacher (teacherno, firstname, surname)
  ```

- `teacherno → firstname, surname` is the only *non-trivial functional depedency* in the relation

- ✅ Each determinant is a candidate key

- ✅ All attribute values are atomic (single values)

- ✅ There is a determinant for each attribute that is not contained in a candidate key

- Thus, *the relation is in BCNF*

Boyce-Codd Normal Form (BCNF)

- Let's consider the following relation:

  ```
  CourseGrade (course_code, studentno, firstname, surname, grade)
  ```

- `studentno → firstname, surname` is one of the *non-trivial functional depedencies* in the relation

- ❌ `studentno` is *not a candidate key* in the relation (so each determinant is *not* a candidate key)

- Thus, *the relation is not in BCNF*

Turning a relation into Boyce-Codd Normal Form

- To convert a *non-BCNF relation to BCNF*, we must decompose the relation in two steps

- Step 1: Find a *non-trivial functional dependency* `X → Y` which violates the BCNF rule (find a determinant that is *not a candidate key*)

- Step 2: Split the original relation in two relations as follows:
  - Create a new relation with all attributes (for example both X and Y) from the dependency. X will be the primary key in the new relation
  - Remove Y attribute(s) from the original relation and leave X in the original relation to act as a foreign key.

- We repeat the steps above until all of our relations are in BCNF

Turning a relation into Boyce-Codd Normal Form

- Let's consider the following relation candidate:

```
CourseOffering (coursecode, offeringno,
                coursename, startdate, teacherno,  surname, firstname)
```

- In the first step, we identify the *non-trivial functional dependencies*:
    - `{coursecode, offeringno} → coursename, startdate, teacherno, surname, firstname`
    - `coursecode → coursename`
    - `teacherno → surname, firstname`

- Then, we identify functional dependencies where the determinant is *not a candidate key*

- There's two such cases: `coursecode → coursename` and `teacherno → surname, firstname`

Turning a relation into Boyce-Codd Normal Form

- In the second step, to solve these two cases we split the original relation two times

- With `coursecode → coursename` we create a new relation Course with attributes `coursecode` and `coursename`

- The determinant, the `coursecode` will be the primary key for the relation. We'll get the following relation:

```
Course (coursecode, coursename)
```

- Finally, we remove the `coursename` from the CourseOffering relation and leave `coursecode` as a foreign key:

```
CourseOffering (coursecode, offeringno,
                startdate, teacherno,  surname, firstname)
  FK (coursecode) REFERENCES Course(coursecode)
```

Turning a relation into Boyce-Codd Normal Form

- We will repeat the same process with `teacherno → surname, firstname` and the final relations are the following:

```
Course (coursecode, coursename)
Teacher (teacherno, surname, firstname)
CourseOffering (coursecode, offeringno, startdate, teacherno)
    FK (coursecode) REFERENCES Course(coursecode)
    FK (teacherno) REFERENCES Teacher(teacherno)
```

- Finally, we *check the decomposed relations*

- In each relation above each determinant is a candidate key and each attribute non-candidate-key attribute has a determinant

- Therefore, the *relations are in BCNF* and we have successfully removed all the undesired redundancy from the design

Summary

- *Database normalisation* is a formal technique of organizing data in a database in a way that *redundancy* and *incosistency* within the data is eliminated

- We analyze a set *normalisation rules* to determine if a relation is in a certain *normal form* (1NF, 2NF, 3NF, BCNF)

- Normalisation rules determine what kind *functional dependencies* the relation can have

- We can turn a non-BCNF relation into BCNF relations by decomposing the relation