

# Aggregate functions and set operators

- During this week we will learn:
  - How to use the aggregate functions `COUNT` , `SUM` , `AVG` , `MIN` and `MAX`
  - How to combine result tables with `UNION` , `EXCEPT` and `INTERSECT` set operators

# Aggregate functions

- Performing some calculation for multiple rows so that the end result is a *single value* is a common query problem
- Example of such query is calculating the count of rows in a certain table
- For example, how can we calculate the number of courses in the `Course` table?
- Functions that perform such operations are referred to as *aggregate functions*

# The COUNT aggregate function

- The `COUNT` aggregate function returns the *total number of rows* that match the specified criteria:

```
-- what's the number of courses in the Course table?  
SELECT COUNT(*) as number_of_courses FROM Course
```

- The result table contains a single row:

number_of_courses
7

# The COUNT aggregate function

- We can also filter the rows the aggregate function operates on using the `WHERE` clause:

```
-- what's the number of courses with more than 3 credits?  
SELECT COUNT(*) as number_of_courses FROM Course  
WHERE credits > 3
```

- The result table contains a single row:

number_of_courses
2

# The COUNT aggregate function

- We can also provide a column name for the `COUNT` aggregate function in which case the function returns the number of the *non-null values* of the given column:

```
-- what's the number of students with an email address?  
SELECT COUNT(email) as number_of_students_with_email  
FROM Student
```

# The SUM aggregate function

- The `SUM` aggregate function returns the *sum of all the values* of a column:

```
-- what's the sum of salaries of female teachers?  
SELECT SUM(salary) as sum_of_salaries FROM Teacher  
WHERE gender = 'F'
```

- The result table contains a single row:

sum_of_salaries
102273.00

# The AVG aggregate function

- The `AVG` aggregate function returns the *average value* of a column:

```
-- what's the average grade from course with code "a730"?  
SELECT AVG(grade) as average_grade FROM CourseGrade  
WHERE course_code = 'a730'
```

- The result table contains a single row:

average_grade
3

# The AVG aggregate function

- Calculating the average includes a division operation, which can produce decimal numbers
- To avoid losing the decimal part of the average, we need to cast integer column values to a `DECIMAL` type:

```
-- multiplying an integer with 1.0 will end up with a DECIMAL type
SELECT AVG(grade * 1.0) as average_grade FROM CourseGrade
WHERE course_code = 'a730'
```

average_grade
3.333333



# The AVG aggregate function

- We can limit the number of decimal places in the result by using casting the result to a `DECIMAL` type with specific precision and scale:

```
-- use scale of 2 in the DECIMAL type to round to two decimals places
SELECT CAST(AVG(grade * 1.0) AS DECIMAL(9, 2)) as average_grade
FROM CourseGrade
WHERE course_code = 'a730'
```

average_grade
3.33

# The MIN aggregate function

- The `MIN` function returns the *smallest value* of a column

```
-- what's the lowest grade from course with code "a730"?  
SELECT MIN(grade) as lowest_grade FROM CourseGrade  
WHERE course_code = 'a730'
```

- The result table contains a single row:

lowest_grade
1

# The MAX aggregate function

- The `MAX` function returns the *largest value* of a column

```
-- what's the highest grade from course with code "a730"?  
SELECT MAX(grade) as highest_grade FROM CourseGrade  
WHERE course_code = 'a730'
```

- The result table contains a single row:

highest_grade
5

# Multiple aggregate functions in a single query

- We can have multiple aggregate functions in the same query:

```
-- what's the highest and lowest grade from course with code "a730"?  
SELECT MAX(grade) as highest_grade, MIN(grade) as lowest_grade FROM CourseGrade  
WHERE course_code = 'a730'
```


- The result table contains a single row with two columns:

highest_grade	lowest_grade
5	1

# Combining results tables with set operators

- We can use the results from *multiple result tables* using the `UNION`, `EXCEPT`, and `INTERSECT` *set operators*
- For example, the `UNION` operator returns *all* the rows from two or more result tables *without duplicate values*:

```
-- What are all the surnames among teachers and students?  
SELECT surname FROM Teacher  
UNION  
SELECT surname FROM Student
```



# The UNION operator

```
SELECT surname FROM Teacher
```

surname
Huhta
Hellerus

```
SELECT surname FROM Student
```

surname
Kokki
Kuikka


```
SELECT surname FROM Teacher  
UNION  
SELECT surname FROM Student
```

surname
Huhta
Hellerus
Kokki
Kuikka

# The EXCEPT operator

- The `EXCEPT` operator returns only the rows from the first result table that are *not included* in the second result table *without duplicate values*:

```
-- What are the campus cities that no student lives in?  
SELECT city FROM Campus  
EXCEPT  
SELECT city FROM Student
```



# The EXCEPT operator

```
SELECT city FROM Campus
```

city
Helsinki
Vantaa

```
SELECT city FROM Student
```

city
Helsinki
Espoo

```
SELECT city FROM Campus  
EXCEPT  
SELECT city FROM Student
```


city
Vantaa



# The INTERSECT operator

- The `INTERSECT` operator returns only the rows that *exist in both result tables without duplicate values*:

```
-- What are the campus cities that have students living in them?  
SELECT city FROM Campus  
INTERSECT  
SELECT city FROM Student
```



# The EXCEPT operator

```
SELECT city FROM Campus
```

city
Helsinki
Vantaa


```
SELECT city FROM Student
```

city
Helsinki
Espoo

```
SELECT city FROM Campus  
INTERSECT  
SELECT city FROM Student
```

city
Helsinki

# The set operators

-  With set operators, the column names and data types of each `SELECT` statement *must match*:

```
-- ✗ first_name column is missing from the latter SELECT statement.  
-- This will lead into an error.  
SELECT surname, first_name FROM Teacher  
UNION  
SELECT surname FROM Student
```

# Summary

- We can perform calculations on multiple rows so that the end result is a single row using *aggregate functions*
- The `COUNT` aggregate function returns the *total number of rows*
- The `SUM` aggregate function returns the *sum of all the values* of a column
- The `AVG` aggregate function returns the *average of values* of a column
- The `MIN` aggregate function returns the *minimum value* of a column
- The `MAX` aggregate function returns the *maximum value* of a column
- The `UNION` , `EXCEPT` and `INTERSECT` *set operators* can be used to combine results of multiple result tables